

PyTorch Image Classification Common Code Template

Kanru Wang

April 2021

Reference

Training

- <https://www.kaggle.com/mekhdigakhramanian/fork-pytorch-efficientnet-baseline-train-amp-a>
- <https://www.kaggle.com/underwearfitting/single-fold-training-of-resnet200d-lb0-965>
- <https://www.kaggle.com/tomehirata/pytorch-training-rfcx-adas-optimizer-resnest>
- <https://www.kaggle.com/haqishen/train-efficientnet-b0-w-36-tiles-256-lb0-87>
- <https://www.kaggle.com/tanulsingh077/pytorch-metric-learning-pipeline-only-images>
- <https://www.kaggle.com/underwearfitting/pytorch-densenet-arcface-validation-training>

Inference

- <https://www.kaggle.com/kanruwang/cassava-ensemble-efficientnet-resnext-private0-901>
- <https://www.kaggle.com/japandata509/ensemble-resnext50-32x4d-efficientnet-0-903>
- <https://www.kaggle.com/raghaw/ensemble-of-2-best-public-notebooks> (single model)
- <https://www.kaggle.com/kneroma/inference-tpu-rfcx-audio-detection-fast> (no TTA)
- <https://www.kaggle.com/haqishen/panda-inference-w-36-tiles-256> (not normal TTA)
- <https://www.kaggle.com/iafoss/panda-concat-tile-pooling-starter-inference>
- <https://www.kaggle.com/parthdhameliya77/nfnet-l0-efficientnet-b5-ensemble-inference> (no TTA)

Ensemble as an alternative to “train on entire dataset”

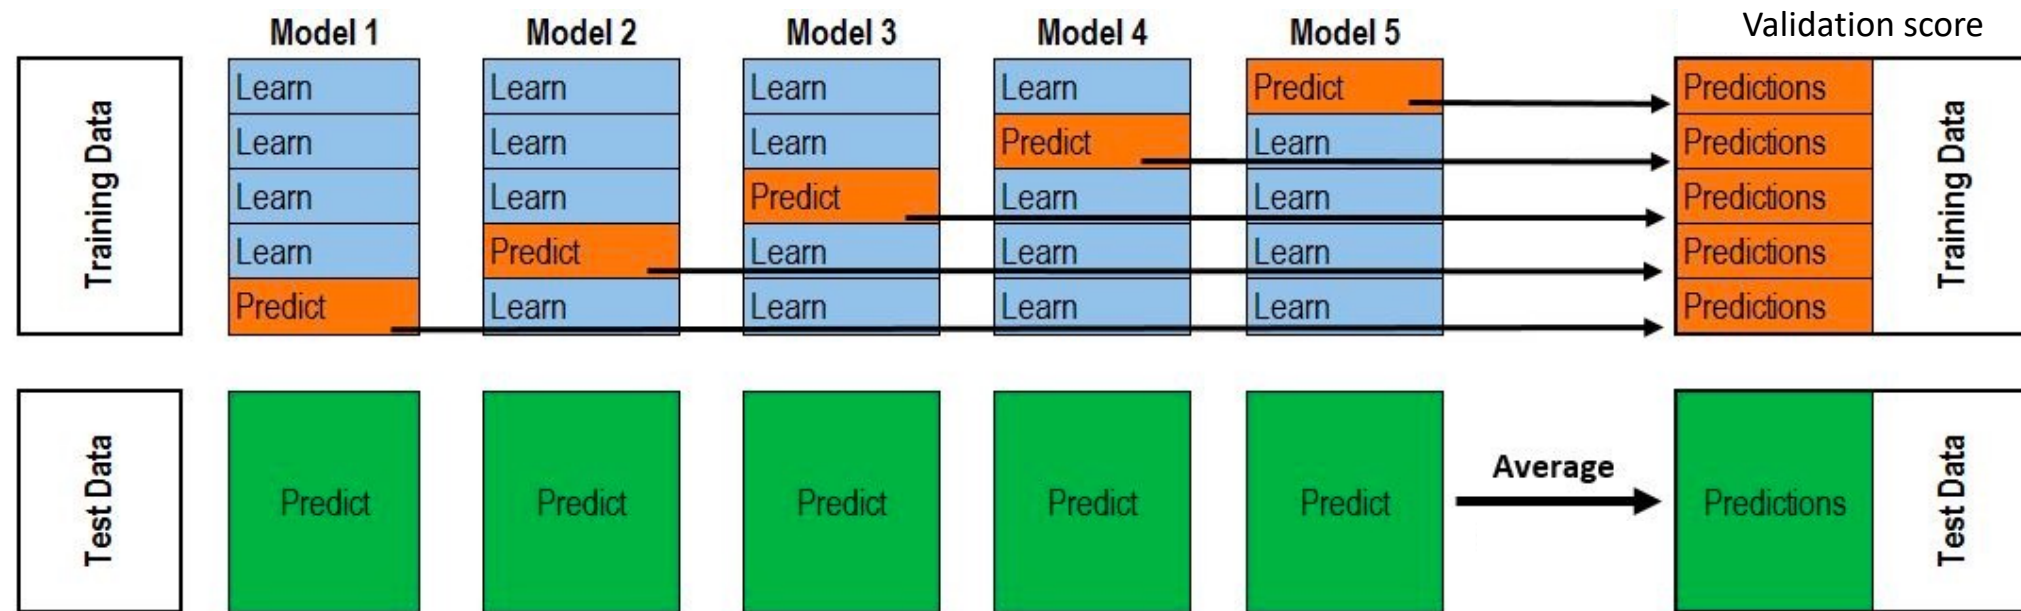
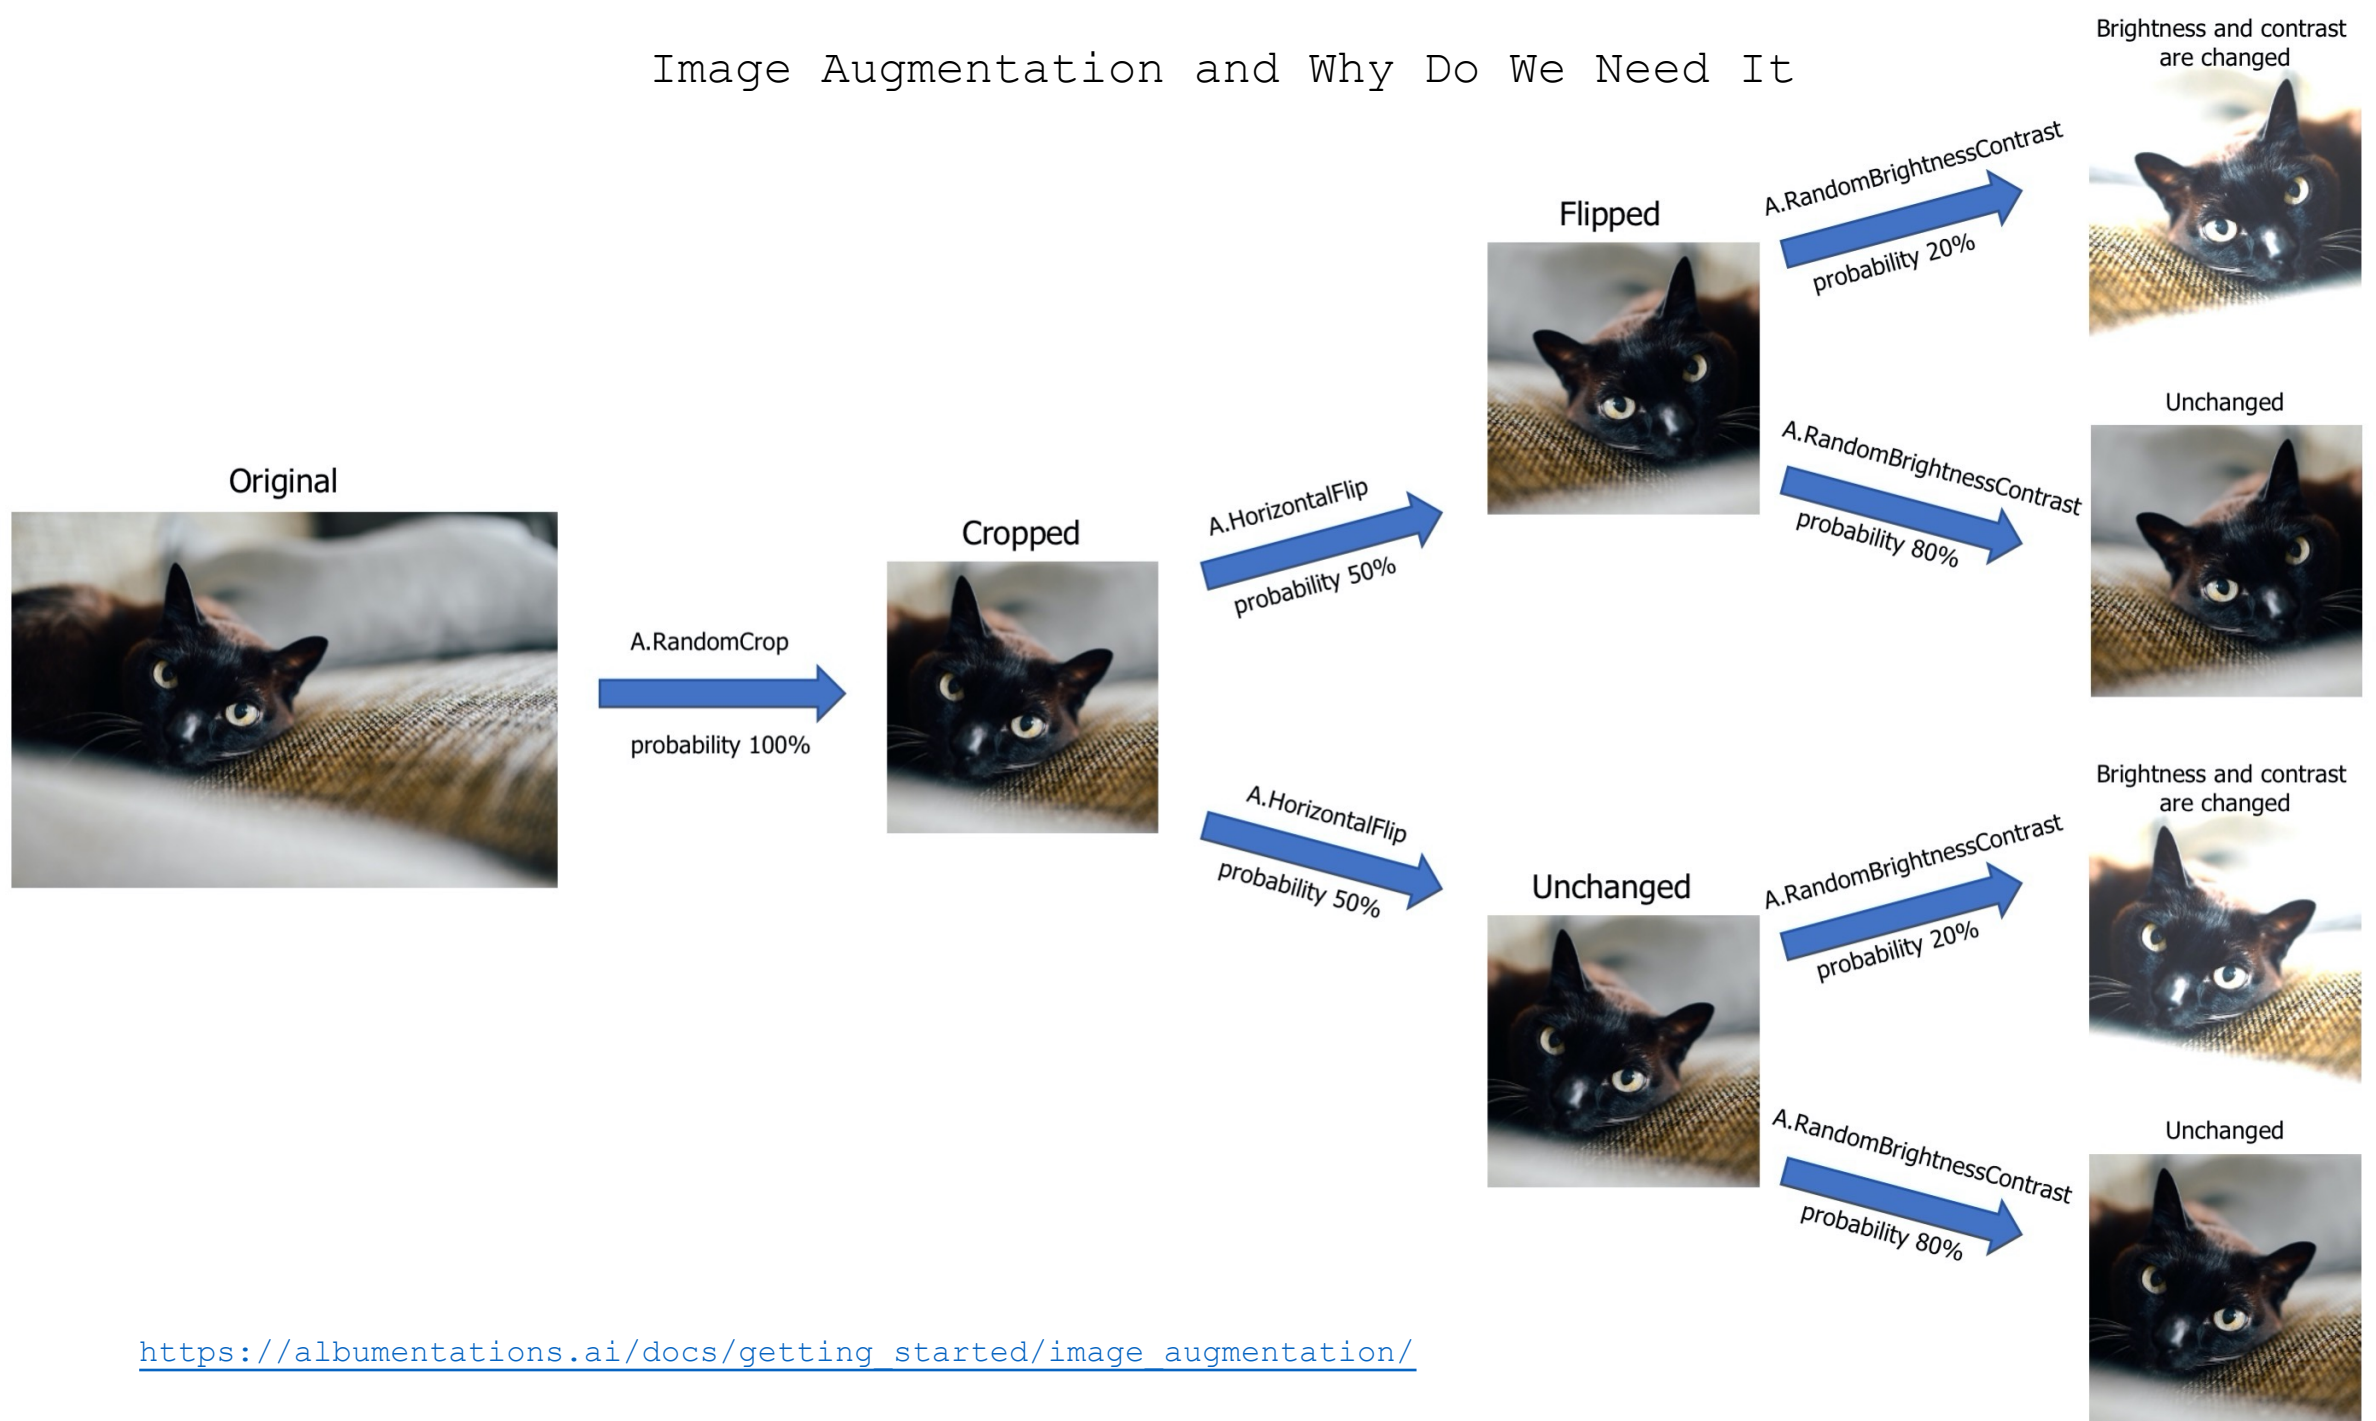
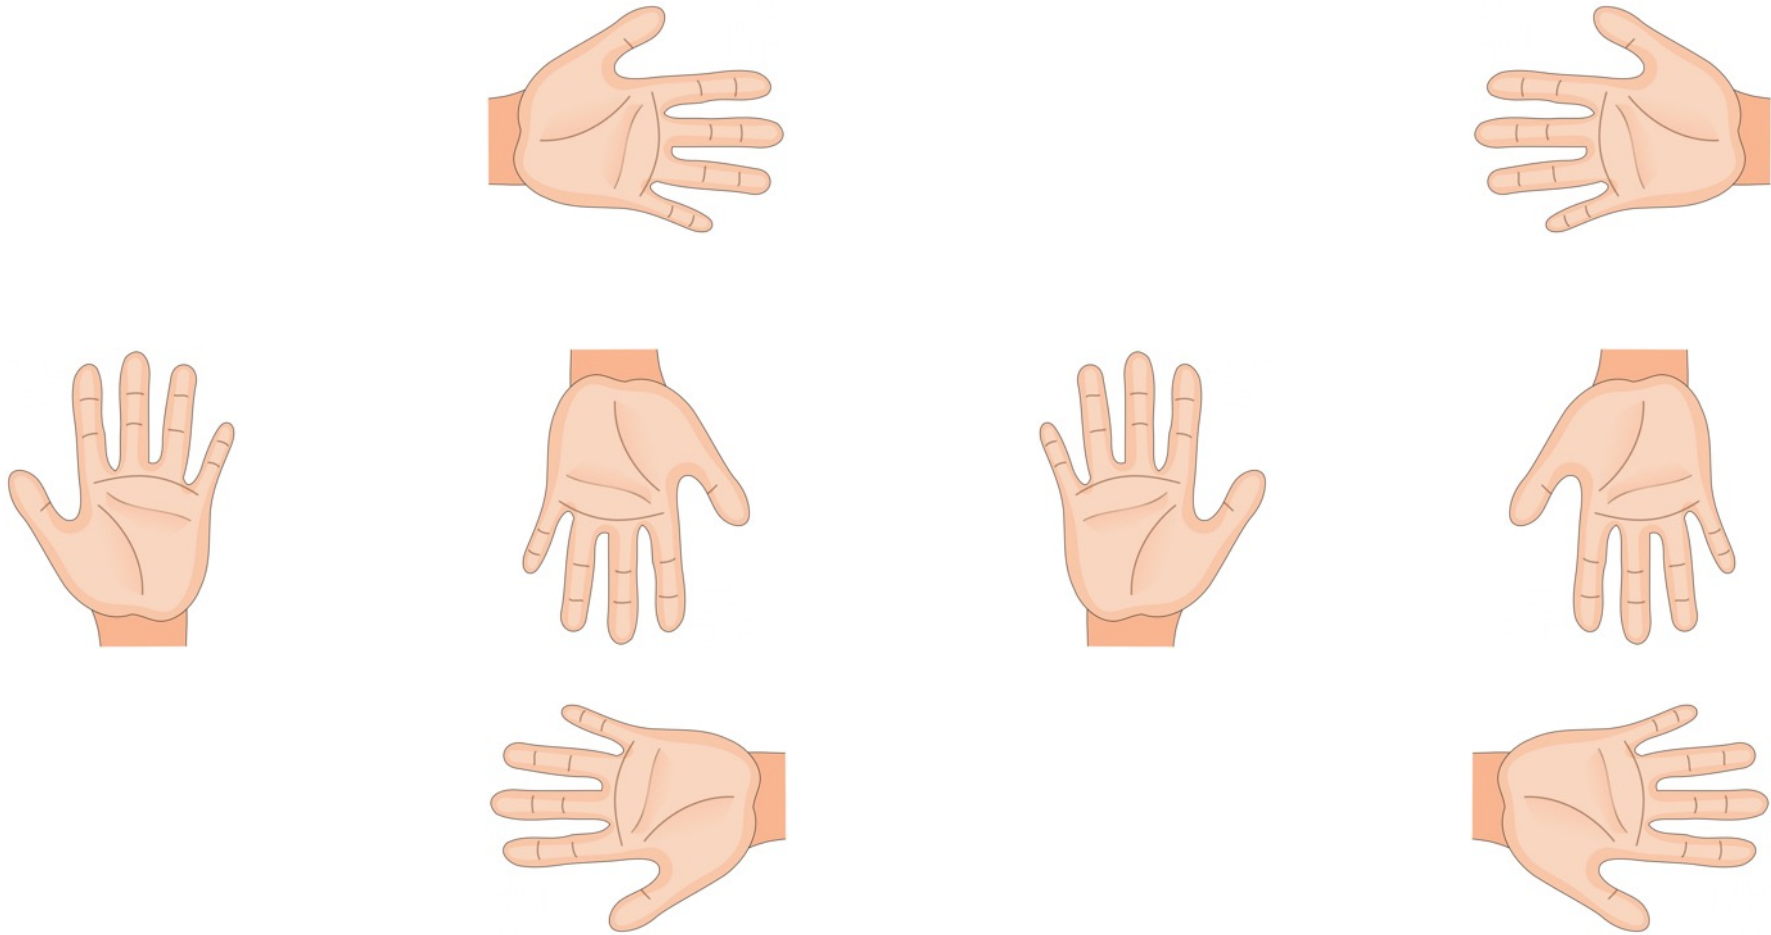


Image Augmentation and Why Do We Need It



8 Test Time Basic Image Augmentation



My understanding:

- During training, do not use the augmentation type that will produce images that are not likely to appear in raw testing/inference images.
- During validation or testing/inference, do not use the augmentation type that has not been used in training.

Training

K folds
(K models)

Train
data
loader

Train
Aug

Val
data
loader

Val
Aug

Epoch

Train

Val

Inference

Test
data loader

Test Aug

K models

TTA

Batch
Pred

```
from albumentations import Compose, HorizontalFlip, VerticalFlip, Transpose, ...

image_augmentation_pipeline_train = Compose([
    RandomResizedCrop
    HorizontalFlip
    OneOf([..., ..., ...])
    ... many more ...
    Normalize
])

image_augmentation_pipeline_valid = Compose([
    ... simpler; just the basic augmentation ...
])

image_augmentation_pipeline_test = Compose([
    ... simpler; just the basic augmentation ...
])
```

Training

K folds
(K models)

Train
data
loader

Train
Aug

Val
data
loader

Val
Aug

Epoch

Train

Val

Inference

Test
dataloader

Test Aug

K models

TTA

Batch
Pred

```
from torch.utils.data import Dataset  
import cv2
```

```
class ImageDataset(Dataset)
```

```
    __init__(...)
```

```
        self.df
```

```
        self.image_augmentation_pipeline
```

```
    __getitem__(self, index:int)
```

```
        row = self.df.loc[index]
```

```
        img = cv2.imread(row["image_path"])
```

```
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB) # OpenCV uses BGR format, Albumentations uses RGB
```

```
        img = self.image_augmentation_pipeline(img) ["image"]
```

```
        if self.output_label == True:
```

```
            return img, row['label'].values
```

```
        else:
```

```
            return img
```

There should be

(1) a folder of images

(2) a csv file that has a column of image file names (or image file path), and a column of labels.

Training

K folds
(K models)

Train
data
loader

Train
Aug

Val
data
loader

Val
Aug

Epoch

Train

Val

Inference

Test
dataloader

Test Aug

K models

TTA

Batch
Pred

```
from torch import nn
```

```
class Model(nn.Module)
```

```
    __init__(...)
```

```
        super().__init__()
```

```
        self.model = model_package.load_pretrained_model("model_type")
```

```
        n_input_features_fully_connected_layer = self.model.fc.in_features
```

```
        self.model.fc = nn.Linear(
```

```
            n_input_features_fully_connected_layer,
```

```
            n_classes_model_output
```

```
        ) # the output dimension of a default model.fc (e.g. 1000) need to be overwritten
```

```
    forward(self, x)
```

```
        " " "
```

```
        forward() is called in nn.Module.__call__(). Should call the module  
        directly using (output = model(data)) instead of model.forward(data).
```

```
        " " "
```

```
        x = self.model(x)
```

```
        return x
```

In PyTorch Image Models (timm), the fully connected layer for EfficientNet and DenseNet is called "self.model.classifier". For ResNet, ResNeXt, SEResNet, and NFNet, it is called "self.model.fc".

Training

K folds
(K models)

Train
data
loader

Train
Aug

Val
data
loader

Val
Aug

Epoch

Train

Val

Inference

Test
dataloader

Test Aug

K models

TTA

Batch
Pred

If need to add Dropout, or change Pooling type, before FC layer

```
class Model(nn.Module)
    __init__(...)
        super().__init__()
        self.model = model_package.load_pretrained_model("model_type")
        n_input_features_fully_connected_layer = self.model.fc.in_features
        self.model.fc = nn.Identity() # set to do nothing, allowing adding Dropout; finally self.fc
        self.fc = nn.Linear(
            n_input_features_fully_connected_layer,
            n_classes_model_output
        )
        self.model.global_pool = nn.Identity() # set to do nothing; can customize pooling
        self.pooling = nn.AdaptiveAvgPool2d(1) # default global_pool is AdaptiveAvgPool2d; we may change
        self.dropout = nn.Dropout(dropout_rate) # (optional)
    forward(self, x)
        x = self.model(x) # result size[batch_size, channels, height, width] e.g. [32, 1536, 16, 16]
        x = self.pooling(x) # result size [batch_size, channels, 1, 1]
        x = x.view(batch_size, -1) # result size [batch_size, channels]
        x = self.dropout(x) # (optional) result size [batch_size, channels]
        x = self.fc(x) # result size [batch_size, n_classes_model_output]
        return x
```

When `self.model(x)` is called, `x` will go through: `backbone(x) -> global_pool[i.e. pool(x) -> flatten(x)] -> fc(x)`

Read `forward_features()`, `create_classifier()`, and `SelectAdaptivePool2d()` in
<https://github.com/rwightman/pytorch-image-models/blob/master/timm/models/resnet.py>

Training

K folds
(K models)

Train
data
loader

Train
Aug

Val
data
loader

Val
Aug

Epoch

Train

Val

Inference

Test
dataloader

Test Aug

K models

TTA

Batch
Pred

```
from torch.utils.data import DataLoader

folds = StratifiedKFold(n_splits, shuffle=True, random_state=42).split(df, df["label"])

for fold, (train_idx, valid_idx) in enumerate(folds):
    df_train = df.loc[train_idx,:].reset_index(drop=True)
    df_valid = df.loc[valid_idx,:].reset_index(drop=True)
    dataset_train = ImageDataset(df_train, image_augmentation_pipeline_train)
    dataset_valid = ImageDataset(df_valid, image_augmentation_pipeline_valid)
    train_loader = DataLoader(dataset_train, train_batch_size, shuffle=True, num_workers=num_workers)
    valid_loader = DataLoader(dataset_valid, valid_batch_size, shuffle=False, num_workers=num_workers)
    model = Model(...).to(device)
    optimizer = torch.optim.Adam(model.parameters(), learning_rate)
    scheduler = ... instantiate the learning rate scheduler ...
    loss_func = ... instantiate the loss function ...
    for epoch in range(num_epochs):
        model = train_one_epoch(...)
        eval_metric = valid_one_epoch(...)
        early_stopping_counter += 1
        if eval_metric > best_eval_metric:
            torch.save(model.state_dict(), path)
            early_stopping_counter = 0
    if early_stopping_counter == early_stopping_limit:
        break
```

Training

K folds
(K models)

Train
data
loader

Train
Aug

Val
data
loader

Val
Aug

Epoch

Train

Val

Inference

Test
dataloader

Test Aug

K models

TTA

Batch
Pred

```
optimizer = torch.optim.Adam(model.parameters(), learning_rate)
device = torch.device("cuda")

def train_one_epoch
    model.train() # for Batchnorm and Dropout to work properly
    for (imgs, labels) in tqdm(train_loader):
        imgs, labels = imgs.to(device), labels.to(device) # moves a model / tensor to GPU
        optimizer.zero_grad() # clears old gradients from the last step's loss.backward()
        logits = model(imgs)
        loss = loss_func(logits, labels) # CrossEntropyLoss or BCEWithLogitsLoss
        loss.backward() # computes the derivative of the loss w.r.t. each parameter using backprop
        optimizer.step() # optimizer updates each parameter using the its gradient and the optimizer rules
        scheduler.step() # can be outside of the for-loop to update LR every epoch, not every batch
        ... attach the loss to a list ...
    ... print the average loss ...

def valid_one_epoch
    model.eval() # for Batchnorm and Dropout to work properly
    with torch.no_grad(): # deactivate backprop to reduce memory usage
        for (imgs, labels) in tqdm(valid_loader):
            imgs, labels = imgs.to(device), labels.to(device)
            logits = model(imgs)
            loss = loss_func(logits, labels)
            ... attach the loss to a list; attach predictions to a list ...
    ... print the average loss; calculate and print the evaluation metric ...
```

Training

K folds
(K models)

Train
data
loader

Train
Aug

Val
data
loader

Val
Aug

Epoch

Train

Val

Inference

Test
dataloader

Test Aug

K models

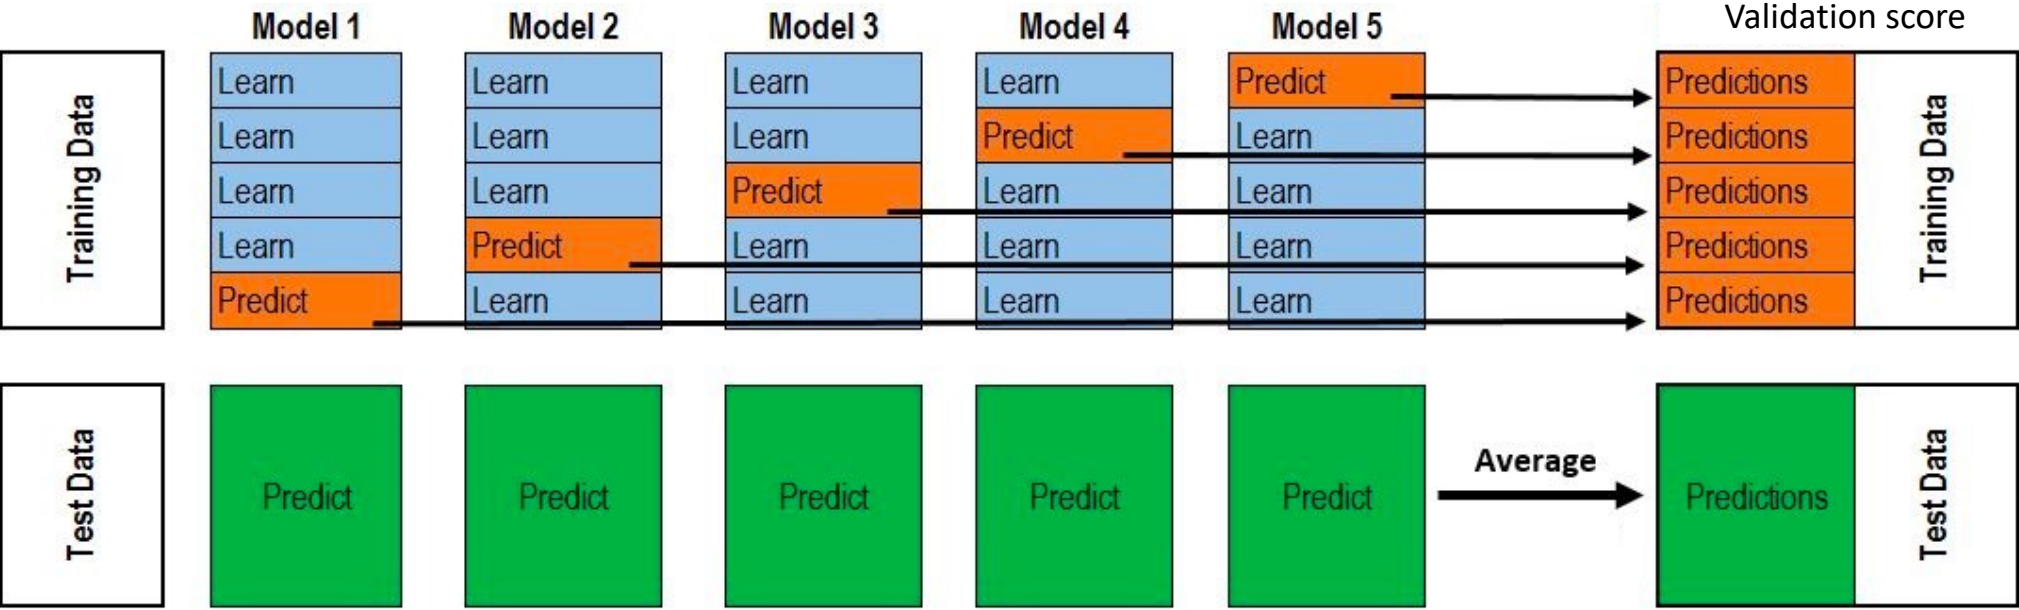
TTA
Batch
Pred

(Optional) Training with Mixed Precision

```
from torch.cuda.amp import autocast, GradScaler # for Mixed Precision
optimizer = torch.optim.Adam(model.parameters(), learning_rate)
device = torch.device("cuda")
scaler = GradScaler()

def train_one_epoch
    model.train()
    for (imgs, labels) in tqdm(train_loader):
        imgs, labels = imgs.to(device), labels.to(device)
        with autocast(): # cast weights from FP32 to FP16 before the forward pass
            logits = model(imgs)
            loss = loss_func(logits, labels) # CrossEntropyLoss or BCEWithLogitsLoss
            scaler.scale(loss).backward() # scale up the loss before the backward pass
            scaler.step(optimizer) # scale down the gradient; then update weights
            scaler.update() # adjust the scaling factor for the next batch
        optimizer.zero_grad() # clears old gradients from the current batch
```

Ensemble as an alternative to “train on entire dataset”



Assume 6 output classes, 4 ensemble models, 8 TTA. For each testing image...

	Model 1								Model 2								Model 3								Model 4								
Class 1																																Avg of 32	The largest value will be the predicted class
Class 2									8 TTA images																							Avg of 32	
Class 3									1	2	3	4	5	6	7	8																Avg of 32	
Class 4																																Avg of 32	
Class 5																																Avg of 32	
Class 6																																Avg of 32	

Training

K folds
(K models)

Train
data
loader

Train
Aug

Val
data
loader

Val
Aug

Epoch

Train

Val

Inference

Test
dataloader

Test Aug

K models

TTA
Batch
Pred

```
dataset_test = ImageDataset(df_test, image_augmentation_pipeline_test_with_flip_and_transpose, output_label=False)

test_loader = DataLoader(dataset_test, test_batch_size, shuffle=False, num_workers=num_workers)

model_paths = ["path_to_fold_1_trained_model", "path_to_fold_2_trained_model" ... ]

Usually exists three for-loops: for each model, for each inference batch, and for each TTA. These loops can be written in any order.

Example 1 (Model Loop -> TTA Loop -> Batch Loop)

logits_epoch_TTA_averaged_all_models = []

for e in range(model_paths):

    model = Model(...).to(device)

    model.load_state_dict(torch.load(model_paths[e]))

    model.eval()

    logits_epoch_TTA_averaged = []

    with torch.no_grad():

        for _ in range(number_of_TTA):

            logits_all = []

            for imgs in tqdm(test_loader):

                imgs = imgs.to(device)

                logits = model(imgs)

                logits_epoch += [torch.softmax(logits, 1).detach().cpu().numpy()]

            logits_epoch_TTA_averaged += [np.concatenate(logits_epoch, axis=0)]

        logits_epoch_TTA_averaged = np.mean(logits_epoch_TTA_averaged, axis=0) # Average over TTA

        logits_epoch_TTA_averaged_all_models.append(logits_epoch_TTA_averaged)

predicted_class = np.mean(logits_epoch_TTA_averaged_all_models, axis=0).argmax(1) # Avg over all models' output; then argmax
```

Training

K folds
(K models)

Train
data
loader

Train
Aug

Val
data
loader

Val
Aug

Epoch

Train

Val

Inference

Test
dataloader

Test Aug

K models

Batch
Pred

TTA

```
dataset_test = ImageDataset(df_test, image_augmentation_pipeline_test_no_flip_or_transpose, output_label=False)
```

```
test_loader = DataLoader(dataset_test, test_batch_size, shuffle=False, num_workers=num_workers)
```

```
model_paths = ["path_to_fold_1_trained_model", "path_to_fold_2_trained_model" ... ]
```

Usually exists three for-loops: for each model, for each inference batch, and for each TTA. These loops can be written in any order.

Example 2 (Model Loop -> Batch Loop -> TTA)

```
logits_TTA_averaged_epoch_all_models = []
```

```
for e in range(model_paths):
```

```
    model = Model(...).to(device)
```

```
    model.load_state_dict(torch.load(model_paths[e]))
```

```
    model.eval()
```

```
    logits_TTA_averaged_epoch = []
```

```
    with torch.no_grad():
```

```
        for imgs in tqdm(test_loader):
```

```
            x = imgs.to(device)
```

```
            x = torch.stack([x, x.flip(-1), x.flip(-2), x.flip(-1, -2), x.transpose(-1, -2), x.transpose(-1, -2).flip(-1),
```

```
                             x.transpose(-1, -2).flip(-2), x.transpose(-1, -2).flip(-1, -2)], 0) # 8 basic TTA
```

```
            x = x.view(-1, 3, img_size, img_size) # result size [inference_batch_size x number_of_TTA, 3, img_size, img_size]
```

```
            logits = model(x)
```

```
            logits_TTA_averaged = logits.view(batch_size, 8, -1).mean(1)
```

```
            logits_TTA_averaged_epoch += [torch.softmax(logits_TTA_averaged, 1).detach().cpu()]
```

```
            logits_TTA_averaged_epoch = torch.cat(logits_TTA_averaged_epoch).cpu().numpy()
```

```
            logits_TTA_averaged_epoch_all_models += [logits_TTA_averaged_epoch]
```

```
predicted_class = np.mean(logits_TTA_averaged_epoch_all_models, axis=0).argmax(1)
```